

07_topic_modelling_falcon

January 22, 2026

1 07_topic_modelling_falcon

This notebook uses an instruction-tuned large language model (Falcon) to extract concise topics from negative Trustpilot reviews, and then applies BERTopic to cluster and summarise those extracted topics. The main steps are:

- Loading the raw filtered review dataset and filtering to negative reviews using rating thresholds defined in the project configuration.
- Optionally sub-sampling the negative reviews for faster experimentation (`SAMPLE_SIZE`), while keeping sampling reproducible via a fixed random seed.
- Running batched topic extraction with the Hugging Face text-generation pipeline using **Falcon-H1-1.5B-Instruct** (`tiiuae/Falcon-H1-1.5B-Instruct`), constrained by:
 - a fixed prompt template
 - maximum input length (`MAX_INPUT_TOKENS`)
 - deterministic decoding (greedy generation with `DO_SAMPLE=False`) for reproducibility
- Preprocessing (using `TextPreprocessor`), parsing and normalising the generated output to extract exactly three short topics per review and filtering out rows with empty or invalid extractions.
- Saving the generated topics and cleaned topic strings as CSV outputs for reporting and downstream modelling.
- Fitting BERTopic (using `BERTopicRunner`) on the cleaned and preprocessed LLM-generated topic strings to identify recurring themes and produce topic-level summaries and visualisations.

The notebook saves outputs to the configured directories:

- Extracted topics and filtered results (CSV): `output/tables/07_topic_modelling_falcon/`
- BERTopic plots (PNG/HTML where applicable): `output/plots/07_topic_modelling_falcon/`
- Saved BERTopic models: `output/models/07_topic_modelling_falcon/`

All experiments were run locally on a MacBook Pro (16-inch, Nov 2024) equipped with an Apple M4 Pro chip (14 CPU cores: 10 performance and 4 efficiency cores) and 24 GB of unified memory, running macOS Tahoe 26.2. Model inference was executed locally; where available, the pipeline can use Apple Silicon acceleration via MPS, otherwise it falls back to CPU. No distributed computing frameworks (e.g. SLURM or OpenPBS) were used. **The notebook has been tested on CPU**

and Apple Silicon (MPS); CUDA support is included for completeness but has not been validated in this project.

The notebook is implemented in Python and relies primarily on pandas, numpy, PyTorch, and Hugging Face Transformers (AutoTokenizer, AutoModelForCausalLM, and pipeline), alongside project-specific utilities for prompt-based topic extraction and BERTopic clustering. For reproducibility, a fixed random seed of **901** was used, and deterministic text generation settings were applied (greedy decoding).

Runtime: On the above hardware, the notebook completes in approximately **4 hours** (MPS) when run on all negative reviews, with runtime dominated by LLM inference (tokenisation and generation). It is recommended **SAMPLE_SIZE = 10** (increase if desired) is set for checking functionality for approximately **1min**.

```
[1]: SAMPLE_SIZE = None # set None to run all rows. Create small sample to check ↴functionality.
```

```
IS_SAMPLE_RUN = SAMPLE_SIZE is not None
SAVE_ARTIFACTS = not IS_SAMPLE_RUN
```

```
[2]: import time
```

```
NOTEBOOK_T0 = time.perf_counter()
print("Notebook timer started.")
```

Notebook timer started.

```
[3]: from pathlib import Path
import sys
import ast
from configparser import ConfigParser

# Resolve project root as the parent of the folder the notebook is currently in
CWD = Path.cwd().resolve()
PROJECT_ROOT = CWD.parent

# Safety fallback
if not (PROJECT_ROOT / "config.ini").exists():
    PROJECT_ROOT = next((p for p in (CWD, *CWD.parents) if (p / "config.ini").exists()), None)
    if PROJECT_ROOT is None:
        raise FileNotFoundError("Could not locate 'config.ini' in the current ↴directory or its parents.")

if str(PROJECT_ROOT) not in sys.path:
    sys.path.insert(0, str(PROJECT_ROOT))

CONFIG = ConfigParser()
CONFIG.read(PROJECT_ROOT / "config.ini")
```

```

print("CONFIG used:")
for section in CONFIG.sections():
    print(f"\n[{section}]")
    for key, value in CONFIG[section].items():
        print(f"{key} = {value}")

```

CONFIG used:

```

[DATA]
data_dir = data
raw_filename = PureGym Customer Reviews.csv
raw_filename_filtered = PureGym Customer Reviews_raw_filtered.csv
preprocessed_filename = PureGym Customer Reviews_preprocessed.csv
preprocessed_filename_sentiment = PureGym Customer
Reviews_preprocessed_sentiment.csv
preprocessed_filename_negative = PureGym Customer
Reviews_preprocessed_negative.csv
preprocessed_filename_non_negative = PureGym Customer
Reviews_preprocessed_non_negative.csv
preprocessed_filename_emotion = PureGym Customer
Reviews_preprocessed_emotion.csv
preprocessed_filename_negative_emotion = PureGym Customer
Reviews_preprocessed_negative_emotion.csv
preprocessed_filename_non_negative_emotion = PureGym Customer
Reviews_preprocessed_non_negative_emotion.csv

[OUTPUT]
plot_dir = output/plots
table_dir = output/tables
model_dir = output/models

[FILTERING]
selected_cols = ["Rating", "Date Experienced", "Review Title", "Review"]
country_code = GB
text_col = Review
detect_language = en
negative_ratings = [1, 2]
emotion_col = Dominant Emotion

[ANALYSIS_DATES]
start_date = 2022-12-17
end_date = 2023-12-17
date_col = Date Experienced

[REPRODUCIBILITY]
seed = 901

```

```

[MODELS]
bertopic_negative = bertopic_negative
bertopic_non_negative = bertopic_non_negative
bertopic_emotion_negative_anger = bertopic_emotion_negative_anger
bertopic_emotion_negative_sadness = bertopic_emotion_negative_sad
bertopic_emotion_negative_joy = bertopic_emotion_negative_joy
bertopic_llm_topics_negative = bertopic_llm_topics_negative
bertopic_negative_reviews_llm_filtered = bertopic_negative_reviews_llm_filtered

[ ]: import pandas as pd
      import numpy as np
      import random
      pd.set_option("display.max_colwidth", None)

      from utils.data_management.data_io import load_csv
      from modelling.bertopic.bertopic_runner import BERTopicRunner
      from utils.processing.text_preprocessor import TextPreprocessor

DATA_DIR = (PROJECT_ROOT / CONFIG["DATA"]["DATA_DIR"])
DATA_DIR.mkdir(parents=True, exist_ok=True)

RAW_FILENAME_FILTERED = CONFIG["DATA"]["RAW_FILENAME_FILTERED"]
RAW_PATH_FILTERED = DATA_DIR / RAW_FILENAME_FILTERED

TEXT_COL = CONFIG["FILTERING"]["TEXT_COL"]
NEGATIVE_RATINGS = ast.literal_eval(CONFIG["FILTERING"].get("NEGATIVE_RATINGS", []
    ↵ "[]"))
SEED = CONFIG["REPRODUCIBILITY"].getint("SEED")

random.seed(SEED)
np.random.seed(SEED)

random.seed(SEED)
np.random.seed(SEED)

TABLE_DIR = PROJECT_ROOT / CONFIG["OUTPUT"]["TABLE_DIR"] /_
    ↵ "07_topic_modelling_falcon"
MODEL_DIR = PROJECT_ROOT / CONFIG["OUTPUT"]["MODEL_DIR"] /_
    ↵ "07_topic_modelling_falcon"
PLOT_DIR = PROJECT_ROOT / CONFIG["OUTPUT"]["PLOT_DIR"] /_
    ↵ "07_topic_modelling_falcon"

TABLE_DIR.mkdir(parents=True, exist_ok=True)
MODEL_DIR.mkdir(parents=True, exist_ok=True)
PLOT_DIR.mkdir(parents=True, exist_ok=True)

```

```
[5]: # Load raw filtered data and filter negative
df_trustpilot = load_csv(str(RAW_PATH_FILTERED))
print("Rows:", len(df_trustpilot))
print("Columns:", df_trustpilot.columns.tolist())

df_negative = df_trustpilot[df_trustpilot[Rating].isin(NEGATIVE_RATINGS)].copy()
print("Negative rows:", len(df_negative))

df_work = df_negative
if SAMPLE_SIZE is not None and len(df_work) > SAMPLE_SIZE:
    df_work = df_work.sample(SAMPLE_SIZE, random_state=SEED)

print("Working rows:", len(df_work))
display(df_work[[TEXT_COL, "Rating"]].head(5))
```

```
Rows: 11300
Columns: ['Rating', 'Date Experienced', 'Review Title', 'Review']
Negative rows: 2368
Working rows: 2368
```

Review \\

10

historically the upkeep of the machines/facilities has been slow but ok. in the last 2-3 months the mens showers have been lukewarm at best. i know other members have noticed this. paying a membership every month for a cold shower is not good enough. let me worry about my carbon footprint. if the cost has gone up increases the membership fee.

16 joined my local puregym when the zero joining fee was on except my local gym was exempt from it and i was charged £15 joining fee. was ready to train over the festive period but i cant do that as my local puregym is closed through out the festive period due to staff shortage... a company thay adverises itself as 24/7 365 days access yet i can not access it when i want because its closed. makes no sence having a policy in place if managers just abuse it. futher more my conversion with the manager team should have stayed between us, our conversion transcript shouldnt have been shared with the staff.

17 almost a week with the gym shut and we still haven't received a single email with an update on how to proceed. i can't go to any other gyms like (finsbury park) being the nearest one - hence why i joined my local in seven sisters and not the other ones. i have to keep coming to check if the gym is re-opened every day because they have absolutely no commitment in keeping us in the loop. there's no maintenance being carried out to try to solve the problem. it's making me think that all this was a catch the £16.99 for the first 6 months offer - and they want us to cancel our membership which has now gone up to £24. 99 a month.

	Rating
10	2
16	1
17	1
21	2
25	1

1.1 Falcon Classes

```
[6]: from __future__ import annotations

import torch
from transformers import AutoModelForCausalLM, AutoTokenizer, pipeline

class FalconInitialiser:
    """Initialise tokenizer, model, and generation pipeline for instruction
    models."""

```

```

@staticmethod
def init_tokenizer(model_name: str) -> AutoTokenizer:
    """Load tokeniser and ensure a pad token is set."""
    tokenizer = AutoTokenizer.from_pretrained(model_name, use_fast=True)

    # Setting pad_token avoids warnings and helps with batched generation. ↴
    # Fall back with eos tokens.
    if tokenizer.pad_token is None:
        tokenizer.pad_token = tokenizer.eos_token

    return tokenizer

@staticmethod
def _has_mps() -> bool:
    """Check whether MPS is built and available."""
    return torch.backends.mps.is_built() and torch.backends.mps.
    ↴is_available()

@staticmethod
def init_model(model_name: str) -> AutoModelForCausalLM:
    """Load model on the best available device."""
    # Prefer CUDA when available
    if torch.cuda.is_available():
        model = AutoModelForCausalLM.from_pretrained(
            model_name,
            dtype=torch.float16,
            device_map="auto",
            low_cpu_mem_usage=True,
        )
        model.eval()
    return model

    # Apple Silicon MPS (next preference)
    if FalconInitialiser._has_mps():
        # Try float16 first for speed/memory, fall back to float32
        try:
            model = AutoModelForCausalLM.from_pretrained(
                model_name,
                dtype=torch.float16,
                low_cpu_mem_usage=True,
            ).to("mps")
        except Exception:
            model = AutoModelForCausalLM.from_pretrained(
                model_name,
                dtype=torch.float32,
                low_cpu_mem_usage=True,
            ).to("mps")

```

```

        model.eval()
        return model

    # Final fallback: CPU
    model = AutoModelForCausallLM.from_pretrained(
        model_name,
        low_cpu_mem_usage=True,
    )
    model.eval()
    return model

@staticmethod
def init_generator(model_name: str):
    """Create a Transformers text-generation pipeline.

    Returns
    ------
    tuple
        (generator, tokenizer) where generator is a text-generation
    ↪pipeline.
    """
    tokenizer = FalconInitialiser.init_tokenizer(model_name)
    model = FalconInitialiser.init_model(model_name)

    # Choose pipeline device argument based on what's available
    # - CUDA: integer GPU index
    # - MPS: "mps"
    # - CPU: -1
    if torch.cuda.is_available():
        device = 0
    elif FalconInitialiser._has_mps():
        device = "mps"
    else:
        device = -1

    # Build a text-generation pipeline for convenience (handles
    ↪tokenisation + generation)
    generator = pipeline(
        task="text-generation",
        model=model,
        tokenizer=tokenizer,
        return_full_text=False,
        device=device,
    )

```

```

        return generator, tokenizer

[ ]: from dataclasses import dataclass
      from typing import Optional

def _prompt_token_len(model_name: str, prompt: str) -> int:
    tokenizer = FalconInitialiser.init_tokenizer(model_name)
    return len(tokenizer.encode(prompt, add_special_tokens=False))

@dataclass(frozen=True)
class FalconParams:
    """Configuration for topic extraction via an instruction model."""

    MODEL_NAME: str = "tiiuae/Falcon-H1-1.5B-Instruct"

    # Batching
    BATCH_SIZE: int = 64

    # Generation controls
    MAX_NEW_TOKENS: int = 30 # response tokens
    DO_SAMPLE: bool = False # greedy decoding for reproducibility

    # Only used when DO_SAMPLE=True
    TEMPERATURE: Optional[float] = None
    TOP_P: Optional[float] = None

    # Output columns
    RAW_COL: str = "Generated Topics"
    EXTRACTED_COL: str = "Extracted Topics"

    # Prompt
    PROMPT: str = (
        "Extract the 3 main topics from this customer review.\n"
        "Rules:\n"
        "-- Return exactly 3 lines and stop.\n"
        "-- Each line must be a short noun phrase (max 6 words).\n"
        "-- No explanation, no extra text, no follow up questions.\n"
        "Format:\n"
        "1. <topic>\n"
        "2. <topic>\n"
        "3. <topic>\n\n"
        "Review:\n"
    )
    # Keep 512 (see diagnostic in cell below) tokens for the review, plus
    # however many tokens the prompt uses.

```

```
MAX_INPUT_TOKENS: int = 512 + _prompt_token_len(MODEL_NAME, PROMPT)
```

```
[8]: MODEL_NAME = FalconParams.MODEL_NAME
generator, tokenizer = FalconInitialiser.init_generator(MODEL_NAME)
```

The fast path is not available because one of `(selective_state_update, causal_conv1d_fn, causal_conv1d_update)` is None. Falling back to the naive implementation. To install follow <https://github.com/state-spaces/mamba/#installation> and <https://github.com/Dao-AI-Lab/causal-conv1d>. Device set to use mps

```
[9]: # Token Length EDA
RUN_TOKEN_LENGTH_DIAGNOSTICS = True
if RUN_TOKEN_LENGTH_DIAGNOSTICS:
    MAXLEN_CANDIDATES = [64, 96, 128, 160, 192, 256, 320, 384, 448, 512]
    df_for_profile = df_negative

    texts = (df_for_profile[TEXT_COL].fillna("").astype(str).str.strip())
    texts = texts.loc[texts.ne("")]

    # Compute token lengths without truncation for each tokeniser
    token_lengths = texts.apply(lambda t: len(tokenizer.encode(t, add_special_tokens=True))).to_numpy()

    # Worst-case length per text across both tokenisers (safe choice for later)
    token_lengths_worst = token_lengths

    # % fully captured for each candidate max length
    captured_pct = [(token_lengths <= L).mean() * 100 for L in
                     MAXLEN_CANDIDATES]
    captured_pct_worst = [(token_lengths_worst <= L).mean() * 100 for L in
                           MAXLEN_CANDIDATES]

    # Summary stats
    percentiles = [50, 75, 90, 95, 97, 99]
    pvals = np.percentile(token_lengths, percentiles)

    p99 = float(np.percentile(token_lengths, 99))

    print("Token length percentiles:")
    for p, v in zip(percentiles, pvals):
        print(f"{p:>2}%%: {int(v)} tokens")

    print(f"\n99th percentile: {int(p99)} tokens")
```

```
    best_L = next((L for L, p in zip(MAXLEN_CANDIDATES, captured_pct_worst) if p >= 99), None)
```

Token length percentiles:

50%: 61 tokens
75%: 114 tokens
90%: 193 tokens
95%: 269 tokens
97%: 315 tokens
99%: 529 tokens

99th percentile: 529 tokens

```
[10]: print("MAX_INPUT_TOKENS:", FalconParams.MAX_INPUT_TOKENS)
```

MAX_INPUT_TOKENS: 595

```
[11]: from __future__ import annotations

from dataclasses import dataclass
from typing import Any, Optional

import pandas as pd
import re
import torch

@dataclass(frozen=True)
class FalconBatchResult:
    """Container for batched generation results."""
    # Raw decoded text returned by the generation pipeline (one per input row)
    raw_outputs: list[str]
    # Parsed topics extracted from the raw output (one per input row)
    extracted_topics: list[str]

class FalconDataCompiler:
    """Apply an instruction model to extract 3 topics per row."""

    def __init__(
        self,
        *,
        df: pd.DataFrame,
        text_col: str,
        generator: Any,
        tokenizer: Any,
        prompt_template: str,
```

```

batch_size: int,
max_input_tokens: int,
max_new_tokens: int,
temperature: Optional[float],
top_p: Optional[float],
do_sample: bool,
raw_output_col: str,
topics_col: str,
) -> None:
    if text_col not in df.columns:
        raise KeyError(f"'{text_col}' not found in DataFrame")

    # Store core inputs
    self._df = df
    self._text_col = text_col

    # Hugging Face text-generation pipeline and matching tokenizer
    self._generator = generator
    self._tokenizer = tokenizer

    # Prompt template and batching controls
    self._prompt_template = str(prompt_template)
    self._batch_size = int(batch_size)

    # Token budgets
    # max_input_tokens controls how many tokens we allow in the full
    # prompt+review
    # max_new_tokens controls how many tokens the model is allowed to
    # generate as output
    self._max_input_tokens = int(max_input_tokens)
    self._max_new_tokens = int(max_new_tokens)

    # Whether to sample (stochastic decoding) or use greedy decoding
    # (deterministic)
    self._do_sample = bool(do_sample)

    # Only keep sampling params when sampling is enabled
    self._temperature: Optional[float] = float(temperature) if (self.
    _do_sample and temperature is not None) else None
        self._top_p: Optional[float] = float(top_p) if (self._do_sample and
    top_p is not None) else None

    # Output column names
    self._raw_output_col = str(raw_output_col)
    self._topics_col = str(topics_col)

```

```

# Support both templates that include "{review}" and templates that do not (experimented both cases in development)
# If "{review}" exists, we split into prefix + suffix so we can token-budget just the review content
if "{review}" in self._prompt_template:
    prefix, suffix = self._prompt_template.split("{review}", 1)
    self._prompt_prefix = prefix
    self._prompt_suffix = suffix
else:
    self._prompt_prefix = self._prompt_template
    self._prompt_suffix = ""

# Pre-tokenise the prefix and suffix once for efficiency and stable budgeting
self._prefix_ids = self._tokenizer.encode(self._prompt_prefix, add_special_tokens=False)
self._suffix_ids = self._tokenizer.encode(self._prompt_suffix, add_special_tokens=False)

# Ensure the prompt scaffold alone does not exceed the model input budget
reserved = len(self._prefix_ids) + len(self._suffix_ids)
if reserved >= self._max_input_tokens:
    raise ValueError(
        "Prompt is longer than MAX_INPUT_TOKENS. Reduce MAX_INPUT_TOKENS or shorten PROMPT."
    )

def apply(self) -> pd.DataFrame:
    """Run generation for all rows and return a copy with new columns."""
    df_out = self._df.copy()

    # Pull the text column into a plain list
    texts = (
        df_out[self._text_col]
        .fillna("")
        .astype(str)
        .tolist()
    )

    # Run batched generation + parsing
    batch_result = self._predict(texts)

    # Append outputs back onto the DataFrame
    df_out[self._raw_output_col] = batch_result.raw_outputs
    df_out[self._topics_col] = batch_result.extracted_topics

```

```

    return df_out

def _predict(self, texts: list[str]) -> FalconBatchResult:
    """Generate in batches and extract numbered topics."""
    # Handle empty inputs
    if not texts:
        return FalconBatchResult(raw_outputs=[], extracted_topics=[])

    # Pre-allocate output lists
    raw_outputs: list[str] = [""] * len(texts)
    extracted_topics: list[str] = [""] * len(texts)

    # Determine pad token for generation; fall back to EOS if pad is unavailable
    pad_id = getattr(self._tokenizer, "pad_token_id", None)
    if pad_id is None:
        pad_id = getattr(self._tokenizer, "eos_token_id", None)

    # Loop over input texts in batches for speed/memory control
    for start in range(0, len(texts), self._batch_size):
        batch = texts[start: start + self._batch_size]

        # Build prompts with per-review truncation to respect max_input_tokens
        prompts = [self._build_prompt(t) for t in batch]

        # Base generation kwargs shared for greedy and sampling modes
        gen_kwargs: dict[str, Any] = dict(
            max_new_tokens=self._max_new_tokens,
            do_sample=self._do_sample,
            pad_token_id=pad_id,
        )

        # Add sampling-specific kwargs only when enabled
        if self._do_sample:
            if self._temperature is not None:
                gen_kwargs["temperature"] = self._temperature
            if self._top_p is not None:
                gen_kwargs["top_p"] = self._top_p

        # Inference only: no gradients needed
        with torch.inference_mode():
            outputs = self._generator(prompts, **gen_kwargs)

        # Normalise the pipeline outputs and extract three numbered topics
        for i, out in enumerate(outputs):

```

```

        idx = start + i
        generated = self._normalise_pipeline_output(out)
        raw_outputs[idx] = generated

        topics = self._extract_numbered_topics(generated, n=3)
        extracted_topics[idx] = " | ".join(topics) if topics else ""

    return FalconBatchResult(raw_outputs=raw_outputs,
                             extracted_topics=extracted_topics)

def _build_prompt(self, review_text: str) -> str:
    """Construct a prompt with review text truncated to fit
    max_input_tokens."""
    review = str(review_text)

    # Compute the token budget available for the review after accounting
    # for prefix + suffix
    budget = self._max_input_tokens - (len(self._prefix_ids) + len(self._suffix_ids))

    # Tokenise the review without special tokens so the budget aligns with
    # prefix/suffix tokenisation
    review_ids = self._tokenizer.encode(review, add_special_tokens=False)

    # Truncate review tokens if needed to ensure the full prompt fits
    # within max_input_tokens
    if len(review_ids) > budget:
        review_ids = review_ids[:budget]

    # Combine and decode back to text for the generation pipeline
    full_ids = self._prefix_ids + review_ids + self._suffix_ids
    return self._tokenizer.decode(full_ids, skip_special_tokens=True)

@staticmethod
def _normalise_pipeline_output(item: Any) -> str:
    """Normalise pipeline outputs to a single generated text string."""
    # HF text-generation pipeline returns:
    # - list[dict] with key "generated_text"
    # - dict with key "generated_text"
    # - other string-like structures depending on wrapper/version
    if isinstance(item, list) and item:
        item0 = item[0]
        if isinstance(item0, dict) and "generated_text" in item0:
            return str(item0["generated_text"])
        return str(item0)

```

```

    if isinstance(item, dict) and "generated_text" in item:
        return str(item["generated_text"])

    return str(item)

@staticmethod
def _extract_numbered_topics(text: str, *, n: int) -> list[str]:
    """Extract up to n topics from numbered output (1., 2., 3.)."""
    s = str(text)

    topics: list[str] = []

    # Preferred: topics appear as separate numbered lines
    for line in s.splitlines():
        match = re.match(r"^\s*([1-9])[\.\.]\s*(.+?)\s*$', line)
        if match:
            topics.append(match.group(2).strip())
        if len(topics) >= n:
            return topics[:n]

    # Fallback: topics might be inline in a single block of text
    inline: list[str] = []
    for m in re.finditer(r"([1-9])[\.\.]\s*(.+?)(?=\\s+[1-9][\.\.])\\s*|$)", s):
        inline.append(m.group(2).strip())
        if len(inline) >= n:
            break

    return inline[:n]

```

```

[12]: compiler = FalconDataCompiler(
    df=df_work,
    text_col=TEXT_COL,
    generator=generator,
    tokenizer=tokenizer,
    prompt_template=FalconParams.PROMPT,
    batch_size=FalconParams.BATCH_SIZE,
    max_input_tokens=FalconParams.MAX_INPUT_TOKENS,
    max_new_tokens=FalconParams.MAX_NEW_TOKENS,
    temperature=FalconParams.TEMPERATURE,
    top_p=FalconParams.TOP_P,
    do_sample=FalconParams.DO_SAMPLE,
    raw_output_col=FalconParams.RAW_COL,
    topics_col=FalconParams.EXTRACTED_COL,
)

df_llm = compiler.apply()

```

```
display(df_llm[[TEXT_COL, FalconParams.RAW_COL, FalconParams.EXTRACTED_COL]].
    head(5))

if SAVE_ARTIFACTS:
    out_all = TABLE_DIR / "llm_topics_negative_processed.csv"
    df_llm.to_csv(out_all, index=False)
    print("Saved:", out_all)
else:
    print("Sample run: skipping save of llm_topics_negative_processed.csv")

df_llm_filtered = df_llm[df_llm[FalconParams.EXTRACTED_COL].astype(str).str.
    strip().ne("")].copy()
print("Rows with extracted topics:", len(df_llm_filtered))

if SAVE_ARTIFACTS:
    out_filtered = TABLE_DIR / "llm_topics_negative_processed_filtered.csv"
    df_llm_filtered.to_csv(out_filtered, index=False)
    print("Saved:", out_filtered)
else:
    print("Sample run: skipping save of llm_topics_negative_processed_filtered.
        csv")

display(df_llm_filtered[[TEXT_COL, FalconParams.EXTRACTED_COL]].head(5))

print("Tokenizer class:", type(tokenizer).__name__)
print("Is fast tokenizer:", getattr(tokenizer, "is_fast", False))
print("Vocab size:", tokenizer.vocab_size)
print("Model max length:", tokenizer.model_max_length)
print("Pad token:", tokenizer.pad_token, tokenizer.pad_token_id)
print("EOS token:", tokenizer.eos_token, tokenizer.eos_token_id)
```

The following generation flags are not valid and may be ignored:
['temperature']. Set `TRANSFORMERS_VERTOSITY=info` for more details.

10
↳ historically the upkeep of the machines/facilities has been slow but ok. in the last 2-3 months the mens showers have been lukewarm at best. i know other members have noticed this. paying a membership every month for a cold shower is not good enough. let me worry about my carbon footprint. if the cost has gone up increases the membership fee.

16 joined my local puregym when the zero joining fee was on except my local gym was exempt from it and i was charged £15 joining fee. was ready to train over the festive period but i cant do that as my local puregym is closed through out the festive period due to staff shortage... a company thay adverisizes itself as 24/7 365 days access yet i can not access it when i want because its closed. makes no sence having a policy in place if managers just abuse it. futher more my conversion with the manager team should have stayed between us, our conversion transcript shouldnt have been shared with the staff.

17 almost a week with the gym shut and we still haven't received a single email ↵ with an update on how to proceed. i can't go to any other gyms like (finsbury ↵ park) being the nearest one - hence why i joined my local in seven sisters and ↵ not the other ones. i have to keep coming to check if the gym is re-opened ↵ every day because they have absolutely no commitment in keeping us in the loop. ↵ there's no maintenance being carried out to try to solve the problem. it's ↵ making me think that all this was a catch the £16.99 for the first 6 months ↵ offer - and they want us to cancel our membership which has now gone up to £24. ↵ 99 a month.

Generated Topics \n\n10 \n\n1. upkeep of machines/facilities\n2.\n\nmens showers\n3. membership fee increase\n\n16 this is a scam. \n\n1. Joining fee issue.\n2. Gym\n\nclosure issue.\n3. Conversion policy issue.\n\n17 i'm not sure if i'll ever get my money back. \n\n1.\ngym shut for almost a week.\n2. i joined\n\n21 no soap. no towels. no hand sanitizer. no hand washing facilities. i am a\n\ndiabetic and i have to clean my own shoes.\n\n25 \n\n1. bad equipment\n\n2. no staff\n3. bad management \n\n\n

Extracted Topics

```
10 upkeep of machines/facilities | mens showers | membership fee increase
16     Joining fee issue. | Gym closure issue. | Conversion policy issue.
17                     gym shut for almost a week. | i joined
21
25                     bad equipment | no staff | bad management

Saved: /Users/Joshua.Dixon/Documents/8_uni/Assignment
2/output/tables/07_topic_modelling_falcon/llm_topics_negative_processed.csv
Rows with extracted topics: 1087
Saved: /Users/Joshua.Dixon/Documents/8_uni/Assignment 2/output/tables/07_topic_m
odelling falcon/llm topics negative processed filtered.csv
```

Review \ 10

historically the upkeep of the machines/facilities has been slow but ok. in the last 2-3 months the mens showers have been lukewarm at best. i know other members have noticed this. paying a membership every month for a cold shower is not good enough. let me worry about my carbon footprint. if the cost has gone up increases the membership fee.

16 joined my local puregym when the zero joining fee
↳ was on except my local gym was exempt from it and i was charged £15 joining
↳ fee. was ready to train over the festive period but i cant do that as my local
↳ puregym is closed through out the festive period due to staff shortage... a
↳ company thay adverises itself as 24/7 365 days access yet i can not access it
↳ when i want because its closed. makes no sence having a policy in place if
↳ managers just abuse it. futher more my conversion with the manager team should
↳ have stayed between us, our conversion transcript shouldnt have been shared
↳ with the staff.

17 almost a week with the gym shut and we still haven't received a single email
↳ with an update on how to proceed. i can't go to any other gyms like (finsbury
↳ park) being the nearest one - hence why i joined my local in seven sisters and
↳ not the other ones. i have to keep coming to check if the gym is re-opened
↳ every day because they have absolutely no commitment in keeping us in the loop.
↳ there's no maintenance being carried out to try to solve the problem. it's
↳ making me think that all this was a catch the £16.99 for the first 6 months
↳ offer - and they want us to cancel our membership which has now gone up to £24.
↳ 99 a month.

25
↳
↳
↳
↳ i dont recommend this gym at all, there a time the boilers was broken down
↳ and they still had it open! no staff never to be seen in this gym! equipment
↳ out dated yet they want to be charging £22.99 as if they are a flag ship gym
↳ in central london! its a standard gym nothing special can see the management
↳ have no interest in the standards of this gym and it definitely not being
↳ managed well.

29
↳
↳
↳ overpriced. its a new gym and very clean but very small. there is a changing
↳ area upstairs shared between everybody. there are about 20 small lockers at
↳ the most and some shelves to put stuff on but you cant lock it. to change you
↳ either need to use toilets or changing cubicles like in shops. there are only
↳ 2 cross trainers and about 6 treadmills. also the gym got closed due to
↳ facilities issues and there isnt any update on when it will reopen so waste of
↳ time and money.

Extracted Topics

10 upkeep of machines/facilities | mens showers | membership fee increase
16 Joining fee issue. | Gym closure issue. | Conversion policy issue.
17 gym shut for almost a week. | i joined
25 bad equipment | no staff | bad management
29 gym size. | facilities issues. | reopen date.

Tokenizer class: PreTrainedTokenizerFast

Is fast tokenizer: True

Vocab size: 65536

Model max length: 100000000000000019884624838656

```
Pad token: <pad> 65536
EOS token: <|end_of_text|> 11
```

```
[13]: # Re-load from disk to ensure downstream steps always use saved artefacts
TOPICS_PATH = PROJECT_ROOT / CONFIG["OUTPUT"]["TABLE_DIR"] /_
↳ "07_topic_modelling_falcon" / "llm_topics_negative_processed.csv"
TOPICS_PATH_FILTERED = PROJECT_ROOT / CONFIG["OUTPUT"]["TABLE_DIR"] /_
↳ "07_topic_modelling_falcon" / "llm_topics_negative_processed_filtered.csv"

df_llm_reloaded = load_csv(TOPICS_PATH)
df_llm_filtered_reloaded = load_csv(TOPICS_PATH_FILTERED)

df_topics_for_bertopic = df_llm_filtered_reloaded.copy()

# Prepare extracted topics for BERTopic
def normalise_llm_topics(text: str) -> str:
    """Normalise extracted topics into a single whitespace-separated string."""
    if not isinstance(text, str):
        return ""
    t = text.strip()
    if not t:
        return ""

    # If any numbering slipped in, remove it
    t = re.sub(r"(?m)^\\s*\\d+\\s*[.\\.]\\-:\\]\\s*", "", t)

    # Convert pipe-separated topics to a single string
    t = t.replace("|", " ")

    # Normalise whitespace
    t = re.sub(r"\\s+", " ", t).strip()
    return t

# This is the Falcon-derived text column for modelling
FALCON_TOPIC_TEXT_COL = "LLM Topics Clean"
df_topics_for_bertopic[FALCON_TOPIC_TEXT_COL] =_
↳ df_topics_for_bertopic[FalconParams.EXTRACTED_COL] ._
↳ apply(normalise_llm_topics)
df_topics_for_bertopic =_
↳ df_topics_for_bertopic[df_topics_for_bertopic[FALCON_TOPIC_TEXT_COL].ne("")]

print("Rows going into BERTopic (topics):", len(df_topics_for_bertopic))
display(df_topics_for_bertopic[[FalconParams.EXTRACTED_COL,_
↳ FALCON_TOPIC_TEXT_COL]].head(5))

STOPWORD_LANGUAGE = "english"
```

```
EXTRA_STOPWORDS = ["pure", "gym", "puregym", # "equipment" retained (removing
    ↵it reduced topic stability and increased fragmentation for this modelling)
    ↵"main", "question", "extracted", "topics", "review", ↵
    ↵"answer" # Additional stopwords associated with generated responses from ↵
    ↵model.
]
PUNCTUATION_PATTERN_TOPIC = r"[-.,\\"'`;:!?()/%]+"
USE_POS_TAGGING_TOPIC = False

df_topics_for_bertopic[FALCON_TOPIC_TEXT_COL] = TextPreprocessor(
    punctuation_pattern=PUNCTUATION_PATTERN_TOPIC,
    extra_stopwords=EXTRA_STOPWORDS,
    use_pos_tagging=USE_POS_TAGGING_TOPIC,
    language=STOPWORD_LANGUAGE,
).transform_many(df_topics_for_bertopic[FALCON_TOPIC_TEXT_COL])

df_topics_for_bertopic = ↵
    ↵df_topics_for_bertopic[df_topics_for_bertopic[FALCON_TOPIC_TEXT_COL]. ↵
    ↵astype(str).str.strip().ne("")]

display(df_topics_for_bertopic[[FalconParams.EXTRACTED_COL, ↵
    ↵FALCON_TOPIC_TEXT_COL]].head(5))
```

Rows going into BERTopic (topics): 1087

Extracted Topics \
0 upkeep of machines/facilities | mens showers | membership fee increase
1 Joining fee issue. | Gym closure issue. | Conversion policy issue.
2 gym shut for almost a week. | i joined
3 bad equipment | no staff | bad management
4 gym size. | facilities issues. | reopen date.

LLM Topics Clean
0 upkeep machine facility men shower membership fee increase

```

1      joining fee issue closure issue conversion policy issue
2                      shut almost week joined
3                      bad equipment staff bad management
4                      size facility issue reopen date

[14]: # Note: SAMPLE_SIZE only affects Falcon extraction. BERTopic is always run on
       ↪the imported (saved) full extracted topic dataset.
# This approach has been adopted as BERTopic struggles to find topics on
       ↪small samples and creates errors when running.

TOP_N_TOPICS = 4
N_WORDS_BARCHART = 5
LABEL = "llm_topics_negative"

UMAP_N_NEIGHBOURS = 15
UMAP_N_COMPONENTS = 5
UMAP_MIN_DIST = 0
UMAP_METRIC = "cosine"

runner = BERTopicRunner(
    model_dir=MODEL_DIR,
    plot_dir=PLOT_DIR,
    table_dir=TABLE_DIR,
    seed=SEED,
    top_n_topics=TOP_N_TOPICS,
    n_words_barchart=N_WORDS_BARCHART,
    min_topic_size=30,
    show_plots=True,
    save_png=True,
    png_scale=2,

    # UMAP controls
    umap_n_neighbors=UMAP_N_NEIGHBOURS,
    umap_n_components=UMAP_N_COMPONENTS,
    umap_min_dist=UMAP_MIN_DIST,
    umap_metric=UMAP_METRIC,
)

result_llm_topics = runner.run(
    df_topics_for_bertopic,
    label=LABEL,
    text_col=FALCON_TOPIC_TEXT_COL,
    verbose=True,
)

print(result_llm_topics.plot_paths)
display(result_llm_topics.topic_info.head(5))

```

```

display(result_llm_topics.top_topics_table.head(5))

2026-01-22 22:39:11,137 - BERTopic - Embedding - Transforming documents to
embeddings.
Batches: 100%|      | 34/34 [00:01<00:00, 29.21it/s]
2026-01-22 22:39:15,155 - BERTopic - Embedding - Completed
2026-01-22 22:39:15,155 - BERTopic - Dimensionality - Fitting the dimensionality
reduction algorithm
2026-01-22 22:39:20,273 - BERTopic - Dimensionality - Completed
2026-01-22 22:39:20,273 - BERTopic - Cluster - Start clustering the reduced
embeddings
2026-01-22 22:39:20,294 - BERTopic - Cluster - Completed
2026-01-22 22:39:20,296 - BERTopic - Representation - Fine-tuning topics using
representation models.
2026-01-22 22:39:20,307 - BERTopic - Representation - Completed
2026-01-22 22:39:20,318 - BERTopic - WARNING: When you use `pickle` to save/load
a BERTopic model, please make sure that the environments in which you save and
load the model are **exactly** the same. The version of BERTopic, its
dependencies, and python need to remain the same.

{'intertopic_distance_html':
PosixPath('/Users/Joshua.Dixon/Documents/8_uni/Assignment 2/output/plots/07_topi
c_modelling_falcon/bertopic_llm_topics_negative_intertopic_distance.html'),
'intertopic_distance_png':
PosixPath('/Users/Joshua.Dixon/Documents/8_uni/Assignment 2/output/plots/07_topi
c_modelling_falcon/bertopic_llm_topics_negative_intertopic_distance.png'),
'barchart_top_topics_html':
PosixPath('/Users/Joshua.Dixon/Documents/8_uni/Assignment 2/output/plots/07_topi
c_modelling_falcon/bertopic_llm_topics_negative_barchart_top4.html'),
'barchart_top_topics_png':
PosixPath('/Users/Joshua.Dixon/Documents/8_uni/Assignment 2/output/plots/07_topi
c_modelling_falcon/bertopic_llm_topics_negative_barchart_top4.png'),
'heatmap_all_topics_html':
PosixPath('/Users/Joshua.Dixon/Documents/8_uni/Assignment 2/output/plots/07_topi
c_modelling_falcon/bertopic_llm_topics_negative_heatmap_all_topics.html'),
'heatmap_all_topics_png':
PosixPath('/Users/Joshua.Dixon/Documents/8_uni/Assignment 2/output/plots/07_topi
c_modelling_falcon/bertopic_llm_topics_negative_heatmap_all_topics.png'),
'heatmap_top_topics_html':
PosixPath('/Users/Joshua.Dixon/Documents/8_uni/Assignment 2/output/plots/07_topi
c_modelling_falcon/bertopic_llm_topics_negative_heatmap_top4.html'),
'heatmap_top_topics_png':
PosixPath('/Users/Joshua.Dixon/Documents/8_uni/Assignment 2/output/plots/07_topi
c_modelling_falcon/bertopic_llm_topics_negative_heatmap_top4.png'),
'topic_info_csv': PosixPath('/Users/Joshua.Dixon/Documents/8_uni/Assignment 2/ou
tput/tables/07_topic_modelling_falcon/bertopic_llm_topics_negative_topic_info.cs
v'), 'top_topics_csv': PosixPath('/Users/Joshua.Dixon/Documents/8_uni/Assignment
2/output/tables/07_topic_modelling_falcon/bertopic_llm_topics_negative_top_topic

```

```

s.csv')}
Topic Count \
0 -1 342 -1_machine_equipment_parking_broken
1 0 248 0_membership_customer_service_fee
2 1 85 1_shower_water_cold_room
3 2 81 2_toilet_smell_smelly_room
4 3 79 3_equipment_space_staff_weight

Representation \
0 [machine, equipment, parking, broken, staff, bad, issue, topic, room, trainer]
1 [membership, customer, service, fee, bad, issue, cancellation, email, access, price]
2 [shower, water, cold, room, locker, hot, temperature, need, warm, facility]
3 [toilet, smell, smelly, room, changing, dirty, broken, urine, management, equipment]
4 [equipment, space, staff, weight, cardio, machine, bad, treadmill, broken, bench]

Representative_Docs
0 [equipment broken changing room dirty need, bad music broken equipment poor, customer service, broken machine poor customer service broken locker topic]
1 [bad customer service, bad customer service, customer service bad membership cancelled]
2 [shower cold staff, shower cold, shower bad]
3 [changing room dirty smell bad cleaning poor, tragic changing room smell toilet, aircon working smell lady changing room toilet working]
4 [space good maintenance bad broken equipment, messy cardio machine space weight area, broken equipment bad staff]

Topic Count \
0 0 248
1 1 85
2 2 81
3 3 79

TopWords
0 membership, customer, service, fee, bad, issue, cancellation, email, access, price
1 shower, water, cold, room, locker, hot, temperature, need, warm, facility

```

```
2 toilet, smell, smelly, room, changing, dirty, broken, urine, management, equipment
3 equipment, space, staff, weight, cardio, machine, bad, treadmill, broken, bench
```

```
[15]: elapsed_s = time.perf_counter() - NOTEBOOK_TO
elapsed_m = elapsed_s / 60
print(f"\nTotal runtime: {elapsed_s:.1f} seconds ({elapsed_m:.2f} minutes)")
```

Total runtime: 12,312.7 seconds (205.21 minutes)